


 ongardie / raft.tla

Branch: master ▾ raft.tla / raft.tla

[Find file](#) [Copy path](#) ongardie Fix AppendEntries to only send one entry at a time

8098acb on 3 Dec 2014

1 contributor

485 lines (433 sloc) 20.1 KB

```
1  ----- MODULE raft -----
2  \* This is the formal specification for the Raft consensus algorithm.
3  \*
4  \* Copyright 2014 Diego Ongaro.
5  \* This work is licensed under the Creative Commons Attribution-4.0
6  \* International License https://creativecommons.org/licenses/by/4.0/
7
8  EXTENDS Naturals, FiniteSets, Sequences, TLC
9
10 \* The set of server IDs
11 CONSTANTS Server  实际上多个servr，但不用自己指定几个，叫什么名字等
12
13 \* The set of requests that can go into the log
14 CONSTANTS Value
15
16 \* Server states.
17 CONSTANTS Follower, Candidate, Leader
18
19 \* A reserved value.
20 CONSTANTS Nil
21
22 \* Message types:
23 CONSTANTS RequestVoteRequest, RequestVoteResponse,
24             AppendEntriesRequest, AppendEntriesResponse
25
26 ----
27 \* Global variables
28
29 \* A bag of records representing requests and responses sent from one server
30 \* to another. TLAPS doesn't support the Bags module, so this is a function
31 \* mapping Message to Nat.
32 VARIABLE messages  TLAPS是啥？ TLA+ book里面提到了Bags模块，是个multi set
33                   这个变量实际上应该是二维map
34
35 \* A history variable used in the proof. This would not be present in an
36 \* implementation.
37 \* Keeps track of successful elections, including the initial logs of the
38 \* leader and voters' logs. Set of functions containing various things about
39 \* successful elections (see BecomeLeader).
40 VARIABLE elections
41
42 \* A history variable used in the proof. This would not be present in an
43 \* implementation.
44 \* Keeps track of every log ever in the system (set of logs).
45 VARIABLE alllogs
46
47 ----
48 \* The following variables are all per server (functions with domain Server).
49
50 \* The server's term number.  这些实际上只有currentTerm, votedFor持久化。其他的在restart后丢失了。参见restart()
51 VARIABLE currentTerm
52 \* The server's state (Follower, Candidate, or Leader).
53 VARIABLE state
54 \* The candidate the server voted for in its current term, or
55 \* Nil if it hasn't voted for any.
56 VARIABLE votedFor
```

```

56  serverVars == <<currentTerm, state, votedFor>>
57
58  \* A Sequence of log entries. The index into this sequence is the index of the
59  \* log entry. Unfortunately, the Sequence module defines Head(s) as the entry
60  \* with index 1, so be careful not to use that!
61  VARIABLE log
62  \* The index of the latest entry in the log the state machine may apply.
63  VARIABLE commitIndex 为什么这个可以不持久化?
64  logVars == <<log, commitIndex>>
65
66  \* The following variables are used only on candidates: 实际上在推选自己当leader期间用
67  \* The set of servers from which the candidate has received a RequestVote
68  \* response in its currentTerm.
69  VARIABLE votesResponded
70  \* The set of servers from which the candidate has received a vote in its
71  \* currentTerm.
72  VARIABLE votesGranted
73  \* A history variable used in the proof. This would not be present in an
74  \* implementation.
75  \* Function from each server that voted for this candidate in its currentTerm
76  \* to that voter's log.
77  VARIABLE voterLog
78  candidateVars == <<votesResponded, votesGranted, voterLog>>
79
80  \* The following variables are used only on leaders: 仅仅leader维护的内存变量
81  \* The next entry to send to each follower.
82  VARIABLE nextIndex
83  \* The latest entry that each follower has acknowledged is the same as the
84  \* leader's. This is used to calculate commitIndex on the leader.
85  VARIABLE matchIndex
86  leaderVars == <<nextIndex, matchIndex, elections>>
87
88  \* End of per server variables.
89  ----
90  \* All variables; used for stuttering (asserting state hasn't changed). 这个是为了表示方便, 比如vars'
91  vars == <<messages, allLogs, serverVars, candidateVars, leaderVars, logVars>> 那些不变的, 必须明确说unchanged
92
93  ----
94  \* Helpers
95
96  \* The set of all quorums. This just calculates simple majorities, but the only
97  \* important property is that every quorum overlaps with every other. 这个定义, 只要server个数>=0, 还是你能表达原意的
98  Quorum == {i \in SUBSET(Server) : Cardinality(i) * 2 > Cardinality(Server)}
99
100
101  \* The term of the last entry in a log, or 0 if the log is empty.
102  LastTerm(xlog) == IF Len(xlog) = 0 THEN 0 ELSE xlog[Len(xlog)].term
103
104  \* Helper for Send and Reply. Given a message m and bag of messages, return a
105  \* new bag of messages with one more m in it.
106  WithMessage(m, msgs) == 这个\in DOMAIN, 与\in不是一码事, 啥含义呢? A function f has a domain, written domain f, and it
107  IF m \in DOMAIN msgs THEN assigns to each element x of its domain the value f [ x ]
108  [msgs EXCEPT ![m] = msgs[m] + 1] 也就是说, msgs是个函数, m 属于这个函数的domain. 实际上, 就是m可以作为下标使用, 或者映射的key使用
109  ELSE
110  msgs @@ (m :> 1)
111
112  \* Helper for Discard and Reply. Given a message m and bag of messages, return
113  \* a new bag of messages with one less m in it.
114  WithoutMessage(m, msgs) ==
115  IF m \in DOMAIN msgs THEN
116  [msgs EXCEPT ![m] = msgs[m] - 1] 这样不会减为0么? 却没有删除掉!
117  ELSE
118  msgs
119
120  \* Add a message to the bag of messages.
121  Send(m) == messages' = WithMessage(m, messages)
122
123  \* Remove a message from the bag of messages. Used when a server is done

```

```

124  \* processing a message.
125  Discard(m) == messages' = WithoutMessage(m, messages)
126
127  \* Combination of Send and Discard
128  Reply(response, request) ==
129      messages' = WithoutMessage(request, WithMessage(response, messages))
130
131  \* Return the minimum value from a set, or undefined if the set is empty.
132  Min(s) == CHOOSE x \in s : \A y \in s : x <= y
133  \* Return the maximum value from a set, or undefined if the set is empty.
134  Max(s) == CHOOSE x \in s : \A y \in s : x >= y
135
136  ----
137  \* Define initial values for all variables
138
139  InitHistoryVars == /\ elections = {}
140                    /\ allLogs  = {}
141                    /\ voterLog = [i \in Server |-> [j \in {} |-> <<>>]]
142  InitServerVars == /\ currentTerm = [i \in Server |-> 1]
143                    /\ state      = [i \in Server |-> Follower]
144                    /\ votedFor   = [i \in Server |-> Nil]
145  InitCandidateVars == /\ votesResponded = [i \in Server |-> {}]
146                       /\ votesGranted  = [i \in Server |-> {}]
147  \* The values nextIndex[i][i] and matchIndex[i][i] are never read, since the
148  \* leader does not send itself messages. It's still easier to include these
149  \* in the functions.
150  InitLeaderVars == /\ nextIndex  = [i \in Server |-> [j \in Server |-> 1]]
151                    /\ matchIndex = [i \in Server |-> [j \in Server |-> 0]]
152  InitLogVars == /\ log          = [i \in Server |-> << >>]
153                 /\ commitIndex = [i \in Server |-> 0]
154  Init == /\ messages = [m \in {} |-> 0]  用{}, 告诉系统, 这是个map?
155          /\ InitHistoryVars
156          /\ InitServerVars
157          /\ InitCandidateVars
158          /\ InitLeaderVars
159          /\ InitLogVars
160
161  ----
162  \* Define state transitions
163
164  \* Server i restarts from stable storage.
165  \* It loses everything but its currentTerm, votedFor, and log.
166  Restart(i) ==
167      /\ state'      = [state EXCEPT ![i] = Follower]
168      /\ votesResponded' = [votesResponded EXCEPT ![i] = {}]
169      /\ votesGranted'  = [votesGranted EXCEPT ![i] = {}]
170      /\ voterLog'     = [voterLog EXCEPT ![i] = [j \in {} |-> <<>>]]
171      /\ nextIndex'    = [nextIndex EXCEPT ![i] = [j \in Server |-> 1]]
172      /\ matchIndex'    = [matchIndex EXCEPT ![i] = [j \in Server |-> 0]]
173      /\ commitIndex'  = [commitIndex EXCEPT ![i] = 0]
174      /\ UNCHANGED <<messages, currentTerm, votedFor, log, elections>>
175
176  \* Server i times out and starts a new election.
177  Timeout(i) == /\ state[i] \in {Follower, Candidate}  Leader, 不会发生超时? 关键是发现也没啥用, 已经是leader
178                /\ state' = [state EXCEPT ![i] = Candidate]
179                /\ currentTerm' = [currentTerm EXCEPT ![i] = currentTerm[i] + 1]
180                \* Most implementations would probably just set the local vote
181                \* atomically, but messaging localhost for it is weaker.
182                /\ votedFor' = [votedFor EXCEPT ![i] = Nil]
183                /\ votesResponded' = [votesResponded EXCEPT ![i] = {}]
184                /\ votesGranted' = [votesGranted EXCEPT ![i] = {}]
185                /\ voterLog' = [voterLog EXCEPT ![i] = [j \in {} |-> <<>>]]
186                /\ UNCHANGED <<messages, leaderVars, logVars>>
187
188  \* Candidate i sends j a RequestVote request.
189  RequestVote(i, j) ==
190      /\ state[i] = Candidate  Candidate状态, 是前提
191      /\ j \notin votesResponded[i]

```

```

192 /\ Send([mtype      |-> RequestVoteRequest,
193         mterm        |-> currentTerm[i],
194         mlastLogTerm  |-> LastTerm(log[i]),
195         mlastLogIndex |-> Len(log[i]),
196         msource       |-> i,
197         mdest         |-> j])
198 /\ UNCHANGED <<serverVars, candidateVars, leaderVars, logVars>>
199
200 \* Leader i sends j an AppendEntries request containing up to 1 entry.
201 \* While implementations may want to send more than 1 at a time, this spec uses
202 \* just 1 because it minimizes atomic regions without loss of generality.
203 AppendEntries(i, j) ==
204   /\ i /= j
205   /\ state[i] = Leader
206   /\ LET prevLogIndex == nextIndex[i][j] - 1
207       prevLogTerm == IF prevLogIndex > 0 THEN
208                       log[i][prevLogIndex].term
209                       ELSE
210                         0
211   \* Send up to 1 entry, constrained by the end of the log.
212   lastEntry == Min({Len(log[i]), nextIndex[i][j]})
213   entries == SubSeq(log[i], nextIndex[i][j], lastEntry)
214   IN Send([mtype      |-> AppendEntriesRequest,
215         mterm         |-> currentTerm[i],
216         mprevLogIndex |-> prevLogIndex,
217         mprevLogTerm  |-> prevLogTerm,
218         mentries      |-> entries,
219         \* mlog is used as a history variable for the proof.
220         \* It would not exist in a real implementation.
221         mlog          |-> log[i],
222         mcommitIndex  |-> Min({commitIndex[i], lastEntry}),
223         msource       |-> i,
224         mdest         |-> j])
225   /\ UNCHANGED <<serverVars, candidateVars, leaderVars, logVars>>
226
227 \* Candidate i transitions to leader.
228 BecomeLeader(i) ==
229   /\ state[i] = Candidate
230   /\ votesGranted[i] \in Quorum
231   /\ state' = [state EXCEPT ![i] = Leader]
232   /\ nextIndex' = [nextIndex EXCEPT ![i] =
233                   [j \in Server |-> Len(log[i]) + 1]]
234   /\ matchIndex' = [matchIndex EXCEPT ![i] =
235                   [j \in Server |-> 0]]
236   /\ elections' = elections \cup
237                   {[eterm |-> currentTerm[i],
238                     eleader |-> i,
239                     elog    |-> log[i],
240                     evotes  |-> votesGranted[i],
241                     evoterLog |-> voterLog[i]]}
242   /\ UNCHANGED <<messages, currentTerm, votedFor, candidateVars, logVars>>
243
244 \* Leader i receives a client request to add v to the log.
245 ClientRequest(i, v) ==
246   /\ state[i] = Leader      如果发的不是leader，直接无视了？不告诉它leader是谁？
247   /\ LET entry == [term |-> currentTerm[i], 里面有term信息
248                   value |-> v]              log', 只有server i的log变了，其他server的log不变
249   newLog == Append(log[i], entry)
250   IN log' = [log EXCEPT ![i] = newLog]      IN的含义: (let def in expr, 只对这个表达式有效)
251   /\ UNCHANGED <<messages, serverVars, candidateVars,
252               leaderVars, commitIndex>>
253
254 \* Leader i advances its commitIndex. leader自身，不含其他
255 \* This is done as a separate step from handling AppendEntries responses,
256 \* in part to minimize atomic regions, and in part so that leaders of
257 \* single-server clusters are able to mark entries committed.
258 AdvanceCommitIndex(i) ==
259   /\ state[i] = Leader

```

```

260  /\ LET \* The set of servers that agree up through index.
261      Agree(index) == {i} \cup {k \in Server :
262                          matchIndex[i][k] >= index}
263  \* The maximum indexes for which a quorum agrees
264  agreeIndexes == {index \in 1..Len(log[i]) :
265                      Agree(index) \in Quorum}
266  \* New value for commitIndex[i]
267  newCommitIndex ==
268      IF /\ agreeIndexes /= {}
269      /\ log[i][Max(agreeIndexes)].term = currentTerm[i]
270      THEN
271          Max(agreeIndexes)
272      ELSE
273          commitIndex[i]  条件不符合，则根本没变
274  IN commitIndex' = [commitIndex EXCEPT ![i] = newCommitIndex]
275  /\ UNCHANGED <<messages, serverVars, candidateVars, leaderVars, log>>
276
277  ----
278  \* Message handlers
279  \* i = recipient, j = sender, m = message
280
281  \* Server i receives a RequestVote request from server j with
282  \* m.mterm <= currentTerm[i].
283  HandleRequestVoteRequest(i, j, m) ==
284      LET logOk == \ / m.mlastLogTerm > LastTerm(log[i])
285                  \ / /\ m.mlastLogTerm = LastTerm(log[i])
286                  /\ m.mlastLogIndex >= Len(log[i])
287      grant == /\ m.mterm = currentTerm[i]
288               /\ logOk
289               /\ votedFor[i] \in {Nil, j}  这个感觉与论文中的反了？==>实际上，UpdateTerm已经修改了server的term，就比mterm大。
290  IN /\ m.mterm <= currentTerm[i]  但是这个感觉比较奇怪，这些/\，是不是一个原子操作？
291      /\ \ / grant /\ votedFor' = [votedFor EXCEPT ![i] = j]  grant为true，执行后面这个。替代了if..else
292      \ / ~grant /\ UNCHANGED votedFor
293      /\ Reply([mtype      |-> RequestVoteResponse,
294                mterm      |-> currentTerm[i],
295                mvoteGranted |-> grant,
296                \* mlog is used just for the `elections' history variable for
297                \* the proof. It would not exist in a real implementation.
298                mlog       |-> log[i],
299                msource     |-> i,
300                mdest       |-> j],
301                m)
302      /\ UNCHANGED <<state, currentTerm, candidateVars, leaderVars, logVars>>
303
304  \* Server i receives a RequestVote response from server j with
305  \* m.mterm = currentTerm[i].
306  HandleRequestVoteResponse(i, j, m) ==
307      \* This tallies votes even when the current state is not Candidate, but
308      \* they won't be looked at, so it doesn't matter.
309      /\ m.mterm = currentTerm[i]
310      /\ votesResponded' = [votesResponded EXCEPT ![i] =
311                          votesResponded[i] \cup {j}]
312      /\ \ / /\ m.mvoteGranted
313          /\ votesGranted' = [votesGranted EXCEPT ![i] =
314                          votesGranted[i] \cup {j}]
315          /\ voterLog' = [voterLog EXCEPT ![i] =
316                          voterLog[i] @@ (j :> m.mlog)]
317      \ / /\ ~m.mvoteGranted
318      /\ UNCHANGED <<votesGranted, voterLog>>
319      /\ Discard(m)
320      /\ UNCHANGED <<serverVars, votedFor, leaderVars, logVars>>
321
322  \* Server i receives an AppendEntries request from server j with
323  \* m.mterm <= currentTerm[i]. This just handles m.entries of length 0 or 1, but
324  \* implementations could safely accept more by treating them the same as
325  \* multiple independent requests of 1 entry.
326  HandleAppendEntriesRequest(i, j, m) ==
327      LET logOk == \ / m.mprevLogIndex = 0

```

```

328         /\ m.mprevLogIndex > 0
329         /\ m.mprevLogIndex <= Len(log[i])
330         /\ m.mprevLogTerm = log[i][m.mprevLogIndex].term
331     IN /\ m.mterm <= currentTerm[i]
332     /\ /\ /\ \* reject request
333         /\ m.mterm < currentTerm[i]
334         /\ m.mterm = currentTerm[i]
335         /\ state[i] = Follower
336         /\ \!not logOk
337     /\ Reply([mtype      |-> AppendEntriesResponse,
338              mterm       |-> currentTerm[i],
339              msuccess     |-> FALSE,
340              mmatchIndex  |-> 0,
341              msource      |-> i,
342              mdest        |-> j],
343              m)
344     /\ UNCHANGED <<serverVars, logVars>>
345 /\ \* return to follower state
346     /\ m.mterm = currentTerm[i]
347     /\ state[i] = Candidate
348     /\ state' = [state EXCEPT ![i] = Follower]
349     /\ UNCHANGED <<currentTerm, votedFor, logVars, messages>>
350 /\ \* accept request
351     /\ m.mterm = currentTerm[i]
352     /\ state[i] = Follower
353     /\ logOk
354     /\ LET index == m.mprevLogIndex + 1
355     IN /\ \* already done with request
356         /\ /\ m.mentries = << >>
357             /\ /\ Len(log[i]) >= index
358                 /\ log[i][index].term = m.mentries[1].term
359                 \* This could make our commitIndex decrease (for
360                 \* example if we process an old, duplicated request),
361                 \* but that doesn't really affect anything.
362                 /\ commitIndex' = [commitIndex EXCEPT ![i] = 根据什么决定修改commitIndex? 而且还是捎带?
363                                     m.mcommitIndex]
364     /\ Reply([mtype      |-> AppendEntriesResponse,
365              mterm       |-> currentTerm[i],
366              msuccess     |-> TRUE,
367              mmatchIndex  |-> m.mprevLogIndex +
368                             Len(m.mentries),
369              msource      |-> i,
370              mdest        |-> j],
371              m)
372     /\ UNCHANGED <<serverVars, logVars>>
373 /\ \* conflict: remove 1 entry
374     /\ m.mentries /= << >>
375     /\ Len(log[i]) >= index
376     /\ log[i][index].term /= m.mentries[1].term
377     /\ LET new == [index2 \in 1..(Len(log[i]) - 1) |->
378                   log[i][index2]]
379     IN log' = [log EXCEPT ![i] = new]
380     /\ UNCHANGED <<serverVars, commitIndex, messages>>
381 /\ \* no conflict: append entry
382     /\ m.mentries /= << >>
383     /\ Len(log[i]) = m.mprevLogIndex
384     /\ log' = [log EXCEPT ![i] =
385                 Append(log[i], m.mentries[1])]
386     /\ UNCHANGED <<serverVars, commitIndex, messages>>
387 /\ UNCHANGED <<candidateVars, leaderVars>>
388
389 \* Server i receives an AppendEntries response from server j with
390 \* m.mterm = currentTerm[i].
391 HandleAppendEntriesResponse(i, j, m) ==
392     /\ m.mterm = currentTerm[i]
393     /\ /\ /\ m.msucces \* successful
394         /\ nextIndex' = [nextIndex EXCEPT ![i][j] = m.mmatchIndex + 1]
395         /\ matchIndex' = [matchIndex EXCEPT ![i][j] = m.mmatchIndex]

```

```

396      \V /\ \not m.msucces \* not successful
397      /\ nextIndex' = [nextIndex EXCEPT ![i][j] =
398          Max({nextIndex[i][j] - 1, 1})]
399      /\ UNCHANGED <<matchIndex>>
400  /\ Discard(m)
401  /\ UNCHANGED <<serverVars, candidateVars, logVars, elections>>
402
403  \* Any RPC with a newer term causes the recipient to advance its term first.
404  UpdateTerm(i, j, m) ==
405      /\ m.mterm > currentTerm[i]
406      /\ currentTerm' = [currentTerm EXCEPT ![i] = m.mterm]
407      /\ state' = [state EXCEPT ![i] = Follower]
408      /\ votedFor' = [votedFor EXCEPT ![i] = Nil]
409      \* messages is unchanged so m can be processed further.
410      /\ UNCHANGED <<messages, candidateVars, leaderVars, logVars>>
411
412  \* Responses with stale terms are ignored.
413  DropStaleResponse(i, j, m) ==
414      /\ m.mterm < currentTerm[i]
415      /\ Discard(m)
416      /\ UNCHANGED <<serverVars, candidateVars, leaderVars, logVars>>
417
418  \* Receive a message.
419  Receive(m) ==
420      LET i == m.mdest
421      j == m.msource
422      IN \* Any RPC with a newer term causes the recipient to advance
423          \* its term first. Responses with stale terms are ignored.
424          \V UpdateTerm(i, j, m)
425          \V /\ m.mtype = RequestVoteRequest
426              /\ HandleRequestVoteRequest(i, j, m)
427          \V /\ m.mtype = RequestVoteResponse
428              /\ \V DropStaleResponse(i, j, m)
429                  \V HandleRequestVoteResponse(i, j, m)
430          \V /\ m.mtype = AppendEntriesRequest
431              /\ HandleAppendEntriesRequest(i, j, m)
432          \V /\ m.mtype = AppendEntriesResponse
433              /\ \V DropStaleResponse(i, j, m)
434                  \V HandleAppendEntriesResponse(i, j, m)
435
436  \* End of message handlers.
437  ----
438  \* Network state transitions
439
440  \* The network duplicates a message
441  DuplicateMessage(m) ==
442      /\ Send(m)
443      /\ UNCHANGED <<serverVars, candidateVars, leaderVars, logVars>>
444
445  \* The network drops a message
446  DropMessage(m) ==
447      /\ Discard(m)
448      /\ UNCHANGED <<serverVars, candidateVars, leaderVars, logVars>>
449
450  ----
451  \* Defines how the variables may transition.
452  Next == /\ \V /\ E i \in Server : Restart(i)
453          \V /\ E i \in Server : Timeout(i)
454          \V /\ E i,j \in Server : RequestVote(i, j)
455          \V /\ E i \in Server : BecomeLeader(i)
456          \V /\ E i \in Server, v \in Value : ClientRequest(i, v)
457          \V /\ E i \in Server : AdvanceCommitIndex(i)
458          \V /\ E i,j \in Server : AppendEntries(i, j)
459          \V /\ E m \in DOMAIN messages : Receive(m)
460          \V /\ E m \in DOMAIN messages : DuplicateMessage(m)
461          \V /\ E m \in DOMAIN messages : DropMessage(m)
462          \* History variable that tracks every log ever:
463          /\ allLogs' = allLogs \cup {log[i] : i \in Server}

```

执行的原子单位是什么？ UpdateTerm会不会执行完了，后面的操作都不执行，中间插入了别的操作？

```
464
465  \* The specification must start with the initial state and transition according
466  \* to Next.
467  Spec == Init /\ [][Next]_vars
468
469  =====
470
471  \* Changelog:
472  \*
473  \* 2014-12-02:
474  \* - Fix AppendEntries to only send one entry at a time, as originally
475  \*   intended. Since SubSeq is inclusive, the upper bound of the range should
476  \*   have been nextIndex, not nextIndex + 1. Thanks to Igor Kovalenko for
477  \*   reporting the issue.
478  \* - Change matchIndex' to matchIndex (without the apostrophe) in
479  \*   AdvanceCommitIndex. This apostrophe was not intentional and perhaps
480  \*   confusing, though it makes no practical difference (matchIndex' equals
481  \*   matchIndex). Thanks to Hugues Evrard for reporting the issue.
482  \*
483  \* 2014-07-06:
484  \* - Version from PhD dissertation
```

//TODO: 这里没有没有让系统保证什么条件一定成立, liveness或者 theorem, 这样运行的意义?
到底在保证什么? 例如paxos例子的voting module中, 有很多。
起码需要保证, 各个server, 同位置的log, 是相同的。